

# ADABOOST

Shingchern D. You

# Ensemble learning

- Stacking
  - ▣ Voting
  - ▣ Training a post classifier
- Bagging
  - ▣ Well known: Random forest
  - ▣ Mentioned before
- Boosting
  - ▣ Well known: AdaBoost (Adaptive Boosting)
  - ▣ Build a strong classifier from many weak classifiers

# Numerical illustration of voting

- Given the following 1-D example (not linearly separable)

X =	0	1	2	3	4	5	6	7	8	9
d =	1	1	1	-1	-1	-1	1	1	1	-1

- 1<sup>st</sup> classifier

X =	0	1	2	3	4	5	6	7	8	9
h1 =	1	1	1	-1	-1	-1	-1	-1	-1	-1

# Numerical illustration of voting

## □ 2<sup>nd</sup> classifier

X =	0	1	2	3	4	5	6	7	8	9
h2 =	1	1	1	1	1	1	1	1	1	-1

## □ 3<sup>rd</sup> classifier

X =	0	1	2	3	4	5	6	7	8	9
h3 =	-1	-1	-1	-1	-1	-1	1	1	1	1

# Numerical illustration of voting

- Perform majority vote

X=	0	1	2	3	4	5	6	7	8	9
h1 =	1	1	1	-1	-1	-1	-1	-1	-1	-1
h2 =	1	1	1	1	1	1	1	1	1	-1
h3 =	-1	-1	-1	-1	-1	-1	1	1	1	1
H=	1	1	1	-1	-1	-1	1	1	1	-1

- All samples are correctly classified

# Numerical illustration of voting

- Although each classifier is a linear weak classifier (i.e., low accuracy), combined classifier is a strong **nonlinear** classifier
- Explain why (**where do we introduce nonlinearity?**)
- AdaBoost follows the same idea, but with weighted sum instead of voting

# AdaBoost algorithm

## □ Symbol definition

- Samples  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^p$
- Desired output  $d_1, \dots, d_n \in \{-1, +1\}$
- Initial weights  $w_{1,1}, \dots, w_{n,1}$  set to  $\frac{1}{n}$  (note: 2nd index is time)
- Weak classifiers  $h: \mathbf{x}_k \rightarrow \{-1, +1\}$

# AdaBoost algorithm

- For  $t = 1 \dots T$ 
  - ▣ Find and save weak classifier  $h_t(\mathbf{x})$  minimize
$$\epsilon_t = \sum_{k=1}^n w_{k,t} \ell(h_t(\mathbf{x}_k) \neq d_k)$$
(Note:  $\epsilon_t$  sometimes could be very small)
  - ▣ Update  $\alpha_t \leftarrow \frac{1}{2} \ln \left( \frac{1-\epsilon_t}{\epsilon_t} \right)$
  - ▣ **Update weights:**  $w_{k,t+1} \leftarrow w_{k,t} \exp(-\alpha_t h_t(\mathbf{x}_k) d_k)$
  - ▣ For  $k = 1 \dots n$  :  $H(\mathbf{x}_k) = \text{sign}(\sum_{z=1}^t \alpha_z h_z(\mathbf{x}_k))$
  - ▣ Stop condition: (1) No error on classifying training data  
Or (2) Upper limit of iterations reached



# AdaBoost algorithm

- Two classifiers in use

- ▣ Use current weak classifier  $h_t(\mathbf{x}_k)$  to update weights

$$w_{k,t+1} \leftarrow w_{k,t} \exp(-\alpha_t \mathbf{h}_t(\mathbf{x}_k) d_k)$$

- ▣ Use combined strong classifier  $H(\mathbf{x}_k)$  to check error samples (but **cannot** be used for weights updating)

$$H(\mathbf{x}_k) = \text{sign} \left( \sum_{z=1}^t \alpha_z h_z(\mathbf{x}_k) \right) == d_k?$$

# AdaBoost algorithm

- For classification after training, use

$$H(\mathbf{x}) = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t(\mathbf{x})\right)$$

- There are several variations on AdaBoost
- The version given here is from **Machine Learning in Action (a good book for engineers)**
- You can compare this algorithm with the one in textbook (original AdaBoost.M1)

# AdaBoost

- AdaBoost has solid theories behind it, but not to be mentioned here
- Some key points in algorithm
  - ▣ Weak classifier (should not be too strong)
  - ▣ Best decision for weighted error in weak classifier
  - ▣ Numerical issues (could be bad)
  - ▣ Very sensitive to noise and outliers

# Weak classifier

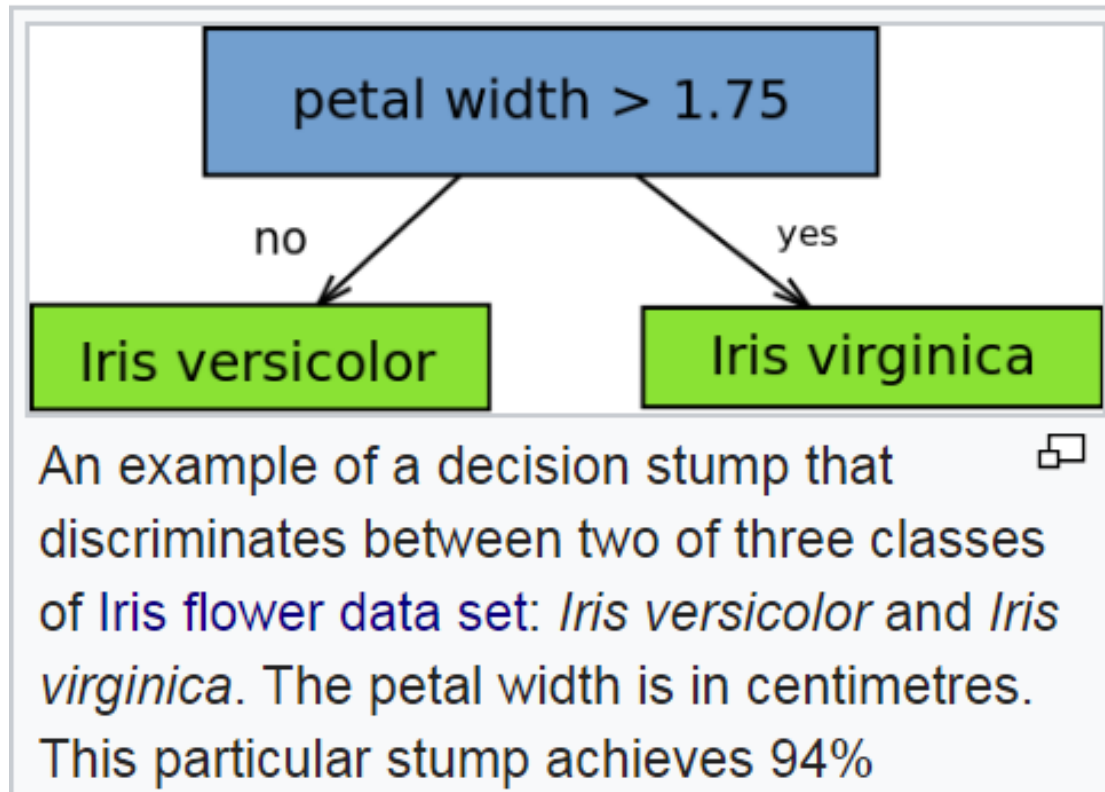
- In the algorithm, we need to search over all possible combinations of parameters to find optimal weighted error

$$\epsilon_t = \sum_{k=1}^n w_{k,t} \ell(h_t(\mathbf{x}_k) \neq y_k)$$

- Not easy with many classifiers (such as SVM)
- One widely used classifier is **decision stump**: making a decision on one feature only

# Weak classifier

- Decision stump example (from wiki)



# Weak classifier

- To find  $\epsilon_t$ , we need to check all possible threshold values for all features
- Consider the following small example:
  - ▣  $P1 = (1, 2.1), C = +1$
  - ▣  $P2 = (2, 1.1), C = +1$
  - ▣  $P3 = (1.3, 1), C = -1$
  - ▣  $P4 = (1, 1), C = -1$
  - ▣  $P5 = (2, 1), C = +1$

# Weak classifier

- For 1<sup>st</sup> feature, we need to check (for example)  
Threshold = {0.9, 1.1, 1.4, 2.1} (other values OK, too)
- For 2<sup>nd</sup> feature, we need to check  
Threshold = { 0.9, 1.05, 1.2, 2.2}
- We also need to know if  $h(x_k) > 0$  means  $C=1$  or  $C=-1$
- Finally, pick the threshold with lowest  $\epsilon_t$

# Weak classifier

- For example, we set  $\text{thr} = 1.4$  in 1<sup>st</sup> feature: if 1<sup>st</sup> feature  $> \text{thr}$ ,  $C = 1$ , else  $C = -1$
- We have only one error in 1<sup>st</sup> iteration ( $t = 1$ )
- Therefore,  $\epsilon_1 = 0.2$ ,  $\alpha_1 = 0.6931$ ,  
 $\mathbf{w}_{\cdot,1} = [0.5, 0.125, 0.125, 0.125, 0.125]^T$
- We can do more steps with same approach



# XOR experiment

- Use 100 samples in XOR as training samples:
- If  $(\text{feature 1}) * (\text{feature 2}) > 0$   
then  $C = 1$ , else  $C = -1$
- Feature 1 and 2 are random numbers
- No error in training set at around 400 iterations (i.e., 400 weak classifiers)

# Using AdaBoost

- **Keep in mind: AdaBoost is very sensitive to noise and outliers (i.e., training samples with wrong classification)**
- Need to use weak classifiers for best performance
- Theories show that AdaBoost also widens the “margin” as SVM does